

04. Randomization

CPSC 535

Kevin A. Wortman



This work is licensed under a Creative Commons Attribution 4.0 International License.

Big Idea: Randomization

Big idea: a *randomized* algorithm deliberately makes random choices

- ▶ con: behavior and/or performance becomes *stochastic* (unpredictable)
- ▶ pro: other aspects can get better (speed, simplicity)
- ▶ often algorithm gets faster/simpler but analysis gets harder
- ▶ (this is a win)

A deterministic algorithm is not randomized.

Recap: Quicksort

Recall:

- ▶ every sorting algorithm takes $\Omega(n \log n)$ time
- ▶ merge sort takes $\Theta(n \log n)$ worst-case time and $\Theta(n)$ space
- ▶ **quicksort synopsis:**
 - ▶ recursive divide-and-conquer
 - ▶ pick random pivot p
 - ▶ rearrange list into two zones: $\leq p$ and $> p$
 - ▶ recurse onto each zone
- ▶ quicksort
 - ▶ is randomized
 - ▶ $\Theta(n \log n)$ expected time
 - ▶ $\Theta(n^2)$ worst-case time
 - ▶ $\Theta(\log n)$ space

Recap: Analyzing Randomized Algorithms

For each randomized algorithm we make two separate efficiency claims:

1. expected time: the focus
2. worst-case time: disclaimer; secondary consideration

“Quicksort takes $\Theta(n \log n)$ expected time and $\Theta(n^2)$ worst-case time.”

Definition of Expected Time

Definition:

algo.'s expected time = $E[\text{algo.'s worst-case time}]$

For random variable X with outcomes x_1, \dots, x_n of probability p_1, \dots, p_n ,

$$E[X] = \sum_{\text{outcome } x_i} p_i x_i$$

So

$$\text{algo.'s expected time} = \sum_{\text{sequence of random choices } i} p_i \cdot (\text{worst-case time steps given } i)$$

The Hiring Problem

```
1: function DETERMINISTIC-HIRE-ASSISTANT( $A$ )
2:    $\text{best} = \text{NIL}$ 
3:   for  $x$  in  $A$  do
4:     if  $\text{best}$  is still  $\text{NIL}$  or  $x$  is better than  $\text{best}$  then
5:        $\text{best} = x$ 
6:     end if
7:   end for
8:   return  $\text{best}$ 
9: end function
```

Analyze the number of **reassignments** (line 5)

Adversarial Analysis

Adversarial analysis: Proof strategy for worst-case analysis

“Adversary”

- ▶ is a fictional opponent character
- ▶ seeks to show algorithm is inefficient
- ▶ has full knowledge of algorithm pseudocode
- ▶ picks least-flattering input

Big idea:

- ▶ randomization makes the adversary's job harder
- ▶ so improves the algorithm's expected performance

Adversarial Analysis of Deterministic Hiring

```
1: function DETERMINISTIC-HIRE-ASSISTANT( $A$ )
2:    $\text{best} = \text{NIL}$ 
3:   for  $x$  in  $A$  do
4:     if  $\text{best}$  is still  $\text{NIL}$  or  $x$  is better than  $\text{best}$  then
5:        $\text{best} = x$ 
6:     end if
7:   end for
8:   return  $\text{best}$ 
9: end function
```

Q: How can an adversary arrange A to maximize reassignments?

Q: What is the worst-case number of reassignments?

Randomized Hiring

```
1: function RANDOMIZED-HIRE-ASSISTANT( $A$ )  
2:   Randomly permute  $A$   
3:    $best = NIL$   
4:   for  $x$  in  $A$  do  
5:     if  $best$  is still  $NIL$  or  $x$  is better than  $best$  then  
6:        $best = x$   
7:     end if  
8:   end for  
9:   return  $best$   
10: end function
```

▷ only change

Observe

- ▶ the same worst-case scenario exists
- ▶ **but**, the adversary cannot force it to occur

Randomized Hiring Analysis

Define

$$X_i = \{1 \text{ if best is reassigned in iteration } i, 0 \text{ otherwise}\}.$$

Observe

$$X_i = 1 \text{ when the } i\text{th element is the maximum so far}$$

and since A is permuted randomly,

$$\Pr\{X_i = 1\} = 1/i \text{ so } E[X_i] = 1/i,$$

and the total number of reassigns is

$$X = 1/1 + 1/2 + 1/3 + 1/4 + \dots + 1/n \in O(\log n).$$


reassignments is $O(\log n)$ expected and $\Theta(n)$ worst-case.

Randomization Patterns

Randomization pattern: approach for using randomization, along with analysis

Best from random order pattern: best only gets reassigned expected $O(\log n)$ times, worst case $\Theta(n)$ times

Balls and Bins



Story to help think about probabilities:

- ▶ b bins that can hold balls
- ▶ throw n balls
- ▶ a ball is equally likely to fall into each bin
- ▶ corresponds to a game called *plinko*

Balls and Bins Q & A

Answers to questions:

- ▶ Q: *After n throws, how many balls does a given bin have?* expected n/b
- ▶ Q: *How many throws before a given bin has a ball?* expected b
- ▶ Q: *How many throws before every bin has a ball?* expected $b \ln b \in \Theta(b \log b)$

Application: Web Server Load Balancer

- ▶ scenario: we have b webserver, n requests coming in, need to route each request to one of the servers $1, \dots, b$
- ▶ adversary could make expensive requests, so if we take turns in a deterministic way, we are vulnerable to a denial-of-service attack
 - ▶ example: round-robin
- ▶ So route requests randomly somehow
- ▶ Idea: choose a **random** server in $\{1, \dots, b\}$
- ▶ Worst-case user experience depends on the heaviest-loaded server

With High Probability (WHP)

Terminology: Let $p(n)$ be the probability that event X occurs as a function of n ; then X occurs *with high probability (w.h.p. or WHP)* when

$$\lim_{n \rightarrow \infty} p(n) = 1.$$

Examples of $p(n)$ bounds considered w.h.p.: $O(1 - \frac{1}{n})$, $O(1 - \frac{1}{n^2})$, $O(1 - \frac{1}{n \log n})$.

Typically, of the form $O(1 - \frac{1}{n^k})$ for $k \geq 1$.

Intuition: for sufficiently large n , event X is practically certain.

Analyzing Random Load Balancing

- ▶ suppose $n = b$
- ▶ suppose the #balls in a bin is its *load*
- ▶ high load is bad
- ▶ *After n throws, what is the maximum load?*

$$\frac{\log n}{\log \log n} + o(1) \text{ w.h.p.}$$

- ▶ Note that for $n \geq 2$,

$$\frac{\log n}{\log \log n} \ll \log n.$$

The Power of Two Random Choices

- ▶ **two random choices**: pick two bins at random, put the ball in *the less-loaded bin*; maximum load becomes

$$\frac{\log \log n}{\log 2} + \Theta(1) \text{ w.h.p.}$$

- ▶ Note that $\frac{\log \log n}{\log 2} \ll \frac{\log n}{\log \log n}$; the choosing dramatically decreases maximum load
- ▶ Intuition: what random events lead to a high load in each scenario?
- ▶ generally, if we make d **random choices**, maximum load is

$$\frac{\log \log n}{\log d} + \Theta(1) \text{ w.h.p.}$$

- ▶ very slow-growing, nearly constant

Load Balancing Patterns

Balance load with one random choice: for n balls in $\Theta(n)$ bins, expected load is $\Theta(1)$ and maximum load is $\Theta\left(\frac{\log n}{\log \log n}\right)$ w.h.p.

Balance load with d random choices: expected load is still $\Theta(1)$, and maximum load is $\Theta\left(\frac{\log \log n}{\log d}\right)$ w.h.p.

Trade-off:

- ▶ one random choice: choosing bin involves only one random number, $\Theta(1)$ time, and does not involve state of bins; *but* load can be more uneven
- ▶ d random choices: choosing bin involves querying the state of $\Theta(d)$ servers, but load is distributed **extremely** evenly

Streaks

- ▶ suppose we flip a fair coin, so $Pr\{\text{heads}\} = Pr\{\text{tails}\} = \frac{1}{2}$
- ▶ *streak*: sequence of the same result (seq. of heads, or seq. of tails)
- ▶ *Q*: After n flips, what is the longest streak?
- ▶ *A*: expected length of the longest streak is $\Theta(\log n)$

Hash Tables

Review hash tables

- ▶ can store a **set** of *keys*
- ▶ or a **map** from keys to arbitrary *values*
- ▶ keys must *hashable*: either integers, or can be mapped deterministically to integers (e.g. strings, floats, tuples of hashable objects, etc.)
- ▶ a search, insert, or delete operation takes $\Theta(1)$ expected time and $\Theta(n)$ worst-case time
- ▶ many variants with trade-offs: chaining vs. open addressing, universal vs. tabular functions, cuckoo, robin hood, etc.

Hash Set Operations

Operation	Pseudocode Example	Exp. Time	W.C. Time
Create empty hash set	<code>S = HashSet()</code>	$\Theta(1)$	$\Theta(1)$
Insert an element	<code>S.insert(x)</code>	$\Theta(1)$	$\Theta(n)$
Remove an element	<code>S.remove(x)</code>	$\Theta(1)$	$\Theta(n)$
Search for an element	<code>if S.contains(x):</code>	$\Theta(1)$	$\Theta(n)$

Recall: A math set cannot have duplicates, so inserting the same x multiple times leaves only one copy in the set.

Hash Map Operations

Operation	Pseudocode Example	Exp. Time	W.C. Time
Create empty hash map	<code>M = HashMap()</code>	$\Theta(1)$	$\Theta(1)$
Insert k, v	<code>M.insert(k, v)</code>	$\Theta(1)$	$\Theta(n)$
Remove key k	<code>M.remove(k)</code>	$\Theta(1)$	$\Theta(n)$
Search for key k	<code>if M.contains(k):</code>	$\Theta(1)$	$\Theta(n)$
Lookup key k	<code>M.get(k)</code>	$\Theta(1)$	$\Theta(n)$

Recall: Each key must be distinct, so re-inserting a new value for key k overwrites the old value.

Reduce-to-Hash-Tables Pattern

- ▶ make critical use of a hash set or hash map
- ▶ replace a $\Theta(n)$ loop with a $\Theta(1)$ expected-time hash table operation
- ▶ good: fast, simple
- ▶ bad: time efficiency becomes *expected*

Duplicate Removal Problem

duplicate removal problem

input: an array $A[1..n]$ of objects

output: a list D of the distinct elements of A (i.e. duplicates are removed)

Duplicate Removal – Baseline Pseudocode

```
1: function REMOVE-DUPPLICATES-BASELINE( $A$ )
2:    $D$  = new list
3:   for  $a$  in  $A$  do
4:     already-present = False
5:     for  $d$  in  $D$  do
6:       if  $a == d$  then
7:         already-present = True
8:       end if
9:     end for
10:    if not already-present then
11:       $D.add(a)$ 
12:    end if
13:  end for
14:  return  $D$ 
15: end function
```

Duplicate Removal – Baseline Analysis

- ▶ outer loop: n iterations
- ▶ inner loop: $\Theta(n)$ time
- ▶ total $\Theta(n^2)$ time
- ▶ bottleneck is the nested loops
- ▶ *Reduce-to-Hash-Tables Pattern*: replace the inner loop with a hash table operation

Duplicate Removal – Improved Algorithm

```
1: function REMOVE-DUPPLICATES-RANDOMIZED(A)
2:   HS = HashSet()
3:   for x in A do
4:     HS.insert(x)
5:   end for
6:   D = insert each element of HS into a list
7:   return D
8: end function
```

▷ match **output:** data type

$\Theta(n)$ expected time, $\Theta(n^2)$ worst-case time.

Duplicate Removal – Analyze Trade-Offs

- ▶ baseline: $\Theta(n^2)$ time; more complicated; not dependent on hash tables knowledge
- ▶ randomized: $\Theta(n)$ expected time, $\Theta(n^2)$ worst-case time; simpler; depends on hash tables
- ▶ randomized is superior
- ▶ pay-off for learning about data structures!

Planning a Hash Map

- ▶ hash set stores *key-value* associations
- ▶ each key is linked to a value
- ▶ *key*: identity of a thing
- ▶ *value*: information associated with that thing
- ▶ insert, remove, search
- ▶ Strategy: fill in the blanks “Given __key__, update __value__ .”
- ▶ Strategy: write out concrete data in a table like

Key	Value
key 1	value 1
key 2	value 2
...	...

Mode Problem

mode problem

input: a list A of n elements

output: a most-frequently-occurring element of A

Note:

- ▶ A is a list, not a set, so A may have duplicates
- ▶ in the event of ties, any of the ties may be output

Mode – Baseline Pseudocode

```
1: function MODE-BASELINE( $A$ )
2:   mode = NIL
3:   mode-count = 0
4:   for  $a$  in  $A$  do
5:     a-count = 0
6:     for  $b$  in  $A$  do
7:       if  $b == a$  then
8:         a-count++
9:       end if
10:    end for
11:    if a-count > mode-count then
12:      mode =  $a$ 
13:      mode-count = a-count
14:    end if
15:  end for
16:  return mode
17: end function
```

Mode – Baseline Analysis

- ▶ total $\Theta(n^2)$ time
- ▶ again, bottleneck is the nested loops
- ▶ *Reduce-to-Hash-Tables Pattern*: replace the inner loop with a hash table operation
- ▶ Pre-compute each element's count
- ▶ Hash map: “Given (an element), update (its count) .”
- ▶ Concrete data :

$$A = \langle 3, 1, 7, 1, 1, 1, 7 \rangle$$

Key	Value
3	1
1	4
7	2

Mode – Second Draft

Hash table is used to speed up first draft, but we haven't initialized it yet.

```
1: function MODE-RANDOMIZED(A)
2:   (somehow create and populate hash map ElementToCount)
3:   mode = NIL
4:   mode-count = 0
5:   for a in A do
6:     a-count = ElementToCount.get(a)
7:     if a-count > mode-count then
8:       mode = a
9:       mode-count = a-count
10:    end if
11:  end for
12:  return mode
13: end function
```

Mode – Final Draft

```
1: function MODE-RANDOMIZED(A)
2:   ElementToCount = HashMap()
3:   for a in A do
4:     if ElementToCount.contains(a) then
5:       ElementToCount.set(a, ElementToCount.get(a) +1)
6:     else
7:       ElementToCount.set(a, 1)
8:     end if
9:   end for
10:  mode = NIL
11:  mode-count = 0
12:  for a in A do
13:    a-count = ElementToCount.get(a)
14:    if a-count > mode-count then
15:      mode = a
16:      mode-count = a-count
17:    end if
18:  end for
19:  return mode
20: end function
```

Mode – Analyze Trade-Offs

- ▶ baseline: $\Theta(n^2)$ time
- ▶ randomized:
 - ▶ first loop: $\Theta(n)$ expected, $\Theta(n^2)$ worst-case
 - ▶ second loop: same
 - ▶ everything else: $\Theta(1)$
 - ▶ total: $\Theta(n)$ expected, $\Theta(n^2)$ worst-case
 - ▶ like the maximum subarray algorithm, refactoring two nested loops into two sequential loops speeds things up
- ▶ again, randomized is superior