

# 06. Linear-Time Sorting and Selection

CPSC 535 ~ Fall 2019

Kevin A. Wortman



CALIFORNIA STATE UNIVERSITY  
**FULLERTON**

September 30, 2019



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

## Big Idea in Computer Science

**“What is old is new again”:** CS principles are solved science; society's needs, economic factors, and fads dictate which are prominent and which are in the background

- ▶ thin clients dominated the mainframe era; thick clients dominated the PC era; thin clients dominate the web app era
- ▶ memory conservation was critical prior to the 90s; programmer labor was more important, until mobile phones
- ▶ Unix rose (70s), fell (80s-90s), rose again (MacOS, iOS, Android, ChromeOS, Linux, PlayStation, embedded)
- ▶ algorithms were considered ivory-tower theory until recently

**Protip:** the CS material that seems irrelevant now, will probably become extremely marketable later in your career

## Big Idea in Algorithm Design

**Parameterized complexity:** algorithm complexity measured both in terms of input size  $n$ , **and** some parameter describing the values in the input

- ▶ machine word size  $W$  (e.g.  $W = 64$  on modern PCs)
- ▶ # distinct values  $k$

**Pseudopolynomial:** polynomial over both  $n$ , and also parameters

- ▶ radix sort takes  $\Theta(nW)$  time
- ▶ strictly speaking  $W$  could be as large as  $n$ , so  $\Theta(nW) = \Theta(n^2)$ , unimpressive
- ▶ in practice all real-world computers have  $W \in \Theta(1)$  so  $\Theta(nW) = \Theta(n)$ , faster than  $\Theta(n \log n)$
- ▶ arguably defying the spirit of the Random Access Model

**Tool to circumvent** lower bounds,  $NP$ -hardness

## The Lower Bound for the Sorting Problem

Recall the precise phrasing of the theorem:

*Any comparison sort algorithm requires  $\Omega(n \log n)$  comparisons in the worst case.*

Bad news:  $O(n \log n)$  “speed limit” for this important problem

Good news:

- ▶ optimal  $\Theta(n \log n)$ -time algorithms: mergesort, heapsort, quicksort
- ▶ loophole: theorem only applies to “comparison sorts”
- ▶ loophole: theorem applies to the general sorting problem, but we could make the problem more specific

## Counting Sort Problem

Recall the classical *sorting problem*:

**input:** a sequence of  $n$  numbers  $A = \langle a_1, a_2, \dots, a_n \rangle$

**output:** a permutation (reordering)  $\langle a'_1, a'_2, \dots, a'_n \rangle$  of the input sequence such that  $a'_1 \leq a'_2 \leq \dots, a'_n$ .

What if the inputs  $a_i$  are all bounded integers?

*counting sort problem:* (changes are underlined)

**input:** an integer  $k \geq 0$ , and a sequence of  $n$  integers  
 $A = \langle a_1, a_2, \dots, a_n \rangle$ , where each  $a_i \in [0, k]$

**output:** same

Turns out this change admits a  $\Theta(n)$ -time algorithm.

## Counting Sort Idea

- ▶ each  $a_i$  can work as an array index
- ▶ count the number of occurrences of each value; create array  $C$  where

$C[x] =$  the number of times that  $x$  appears in  $A$

- ▶ let  $B$  be the output array
- ▶ use the counts in  $C$  to plan out which indices of  $B$  will hold each  $a_i$
- ▶ fill in  $B$  using this information

## Counting Sort Pseudocode

```

1: function COUNTING-SORT( $A, B, k$ )
2:   allocate new array  $C[0, \dots, k]$ , initialized to all zeroes
3:   for  $j$  from 1 to  $A.length$  do
4:      $C[A[j]] + +$   $\triangleright C[i]$  is the number of elements  $= i$ 
5:   end for
6:   for  $i$  from 1 to  $k$  do
7:      $C[i] + = C[i - 1]$   $\triangleright C[i]$  is the number of elements  $\leq i$ 
8:   end for
9:   for  $j$  from  $A.length$  down to 1 do
10:     $B[C[A[j]]] = A[j]$ 
11:     $C[A[j]] - -$ 
12:  end for
13: end function

```

## Counting Sort Analysis

- ▶ allocate  $C$ :  $\Theta(k)$  time
- ▶ first **for** loop:  $\Theta(n)$  time
- ▶ second **for** loop:  $\Theta(k)$  time
- ▶ third **for** loop:  $\Theta(n)$  time
- ▶ total =  $\Theta(k + n + k + n) = \Theta(2k + 2n) = \Theta(k + n)$  time



## When Counting Sort Wins

If  $k \in O(n)$ , then counting sort takes  $\Theta(k + n) = \Theta((n) + n) = \Theta(n)$  time.

Applications where  $k \ll n$

- ▶ DNA sequences have only  $k = 4$  bases A, C, G, T; human genome has  $n \approx 3$  billion bases
- ▶  $n$ -character ASCII string has only  $k = 127$  character values
- ▶ log of website analytics has  $n$  hits but only  $k$  distinct web page URLs; each page is visited many times so  $n \gg k$

## When Counting Sort Loses

In the general sorting problem, each  $a_i$  is unbounded, so the maximum element could use  $\Theta(n)$  bits, have value

$$a_i = 2^n = k$$

and force counting sort to take

$$\Theta(k + n) = \Theta((2^n) + n) = \Theta(2^n)$$

time, which is **exponential** in  $n$  and much more expensive than  $\Theta(n \log n)$

$\implies$  counting sort is optimal when the software designer knows that the input is always a set of  $k$  integers with  $k \in O(n)$

$\implies$  **but** if that is not guaranteed, comparison sorts are still optimal

## Stable Sorting

**stable** sorting algorithm: does not swap order of ties

if  $a_i = a_j$  and  $i < j$  then  $i' < j'$

Ex.: suppose we sort

$[(47, a), (13, c), (28, b), (13, d)]$

by the first element of each pair; a *stable* sort guarantees  $(13, c)$  comes before  $(13, d)$

Stability is a convenient, desirable property

Stable: insertion sort, mergesort, counting sort

Not stable: heapsort, quicksort

# “Radix”

Vocabulary quiz:

- ▶ what does **radix** mean?
- ▶ where else do we use the word “radix”?

## Radix Sort Overview

- ▶ make counting sort more robust to large elements
- ▶ sort one digit at a time
- ▶ i.e. sort by least-significant-digit, then by second-least-significant-digit, ..., sort by most-significant digit
- ▶ e.g. to sort names in a spreadsheet: sort by first name, then by last name
- ▶ originally **used by pre-digital punchcard sorting machines** (what's old...)
- ▶ now used for parallel sort in GPU (...is new again)

## Radix Sort Worked Example

(Sorting one base-10 digit at a time.)

7	7	5	7	9	1	5	2	1	0	5	3
0	5	3	5	2	1	2	2	5	2	2	5
7	6	2	7	6	2	3	3	4	3	3	4
7	9	1	0	5	3	0	5	3	5	2	1
6	7	4	6	7	4	7	6	2	6	7	4
5	2	1	3	3	4	6	7	4	7	6	2
3	3	4	7	7	5	7	7	5	7	7	5
2	2	5	2	2	5	7	9	1	7	9	1

## Radix Sort, 1 bit at a time

```
1: function RADIX-SORT-1( $A$ ,  $W$ )
2:   for  $i$  from 1 to  $W$  do
3:     use a stable sort to sort  $A$  based only on bit position  $i$ 
4:   end for
5: end function
```

Using counting sort as the stable sort, we have  $k = 2$  (bit values 0 or 1) so each loop iteration takes  $\Theta(k + n) = \Theta(2 + n) = \Theta(n)$  time

Clearly  $W$  iterations  $\implies \Theta(nW)$  total time

## Radix Sort, 8 bits at a time

```
1: function RADIX-SORT-8(A, W)
2:   for  $i$  from 1 to  $\lceil W/8 \rceil$  do
3:     stably sort  $A$  on bits  $8i - 7$  through  $8i$ 
4:   end for
5: end function
```

$$k = 2^8 = 256 \in \Theta(1)$$

number of iterations is  $\lceil nW/8 \rceil \in \Theta(n)$

$\implies$  still  $\Theta(nW)$  time, but with different constant factors

Let  $r = \#$ bits per pass; optimal choice of  $r$  minimizes

$$\lceil W/r \rceil \cdot (2^r + n)$$



## Minimum and Maximum

```
1: function MINIMUM(A)
2:   min = A[1]
3:   for  $i$  from 2 to  $A.length$  do
4:     if  $min > A[i]$  then
5:       min = A[i]
6:     end if
7:   end for
8:   return min
9: end function
```

$\Theta(n)$  time

can also find maximum in  $\Theta(n)$  time, or both in  $\Theta(n)$  time

## Selection Problem and Baseline Algorithm

**input:** array of  $n$  numbers  $A = \langle a_1, a_2, \dots, a_n \rangle$ ; index  $i \in \{1, 2, \dots, n\}$

**output:** the  $i$ th smallest element of  $A$

- 1: **function** SELECTION-BY-SORTING( $A, i$ )
- 2:     **return**  $MERGE - SORT(A)[i]$
- 3: **end function**

Clearly  $\Theta(n \log n)$  time

**Surprise:** selection can be solved in only  $\Theta(n)$  time

## Randomized Quicksort Review

```

1: function RQSORT(A, p, r)
2:   if  $p < r$  then
3:      $q = \text{RPART}(A, p, r)$ 
4:     RQSORT(A, p,  $q - 1$ )
5:     RQSORT(a,  $q + 1$ , r)
6:   end if
7: end function

```

Non-stable sort in  $\Theta(n \log n)$   
 expected time but  $\Theta(n^2)$   
 worst-case time

```

1: function RPART(A, p, r)
2:    $i = \text{random in } [p, r]$ 
3:   swap(A[i], A[r])
4:   pivot = A[r]
5:    $i = p$ 
6:   for  $j$  from  $p$  to  $r - 1$  do
7:     if A[j] < pivot then
8:       swap(A[i], A[j])
9:        $i++$ 
10:    end if
11:  end for
12:  swap(A[i], A[r])
13:  return  $i$ 
14: end function

```

## Randomized Selection Overview

- ▶ combining ideas from binary search and quicksort
- ▶ recursively search for  $i$ th smallest element
- ▶ do randomized partition; then
- ▶ three cases
  - ▶ pivot happens to be  $i$ th smallest
  - ▶ need to keep searching before pivot
  - ▶ need to keep searching after pivot
- ▶ expected runtime is  $T(n) \approx T(n/2) + \Theta(n)$
- ▶ counterintuitively, that solves to  $\Theta(n)$

## Randomized Selection Pseudocode

```
1: function RSELECT( $A, p, r, i$ )
2:   if  $p == r$  then
3:     return  $A[p]$                                 ▷ base case, done
4:   end if
5:    $q = \text{RPART}(A, p, r)$                         ▷ partition,  $q$  is pivot index
6:    $k = q - p + 1$                                 ▷  $k$  = number of elements before pivot
7:   if  $i == k$  then
8:     return  $A[q]$                                 ▷ pivot is answer
9:   else if  $i < k$  then
10:    return  $\text{RSELECT}(A, p, q - 1, i)$ 
11:  else
12:    return  $\text{RSELECT}(A, q + 1, r, i - k)$  ▷  $i$  decreases by  $k$ 
13:  end if
14: end function
```

## Randomized Selection Analysis

- ▶ at most one recursive call, on  $n/2$  elements on average
- ▶ partitioning takes  $\Theta(n)$  time
- ▶ rest of algorithm takes  $\Theta(1)$  time
- ▶ expected running time

$$T(n) = T(n/2) + \Theta(n)$$

which is only  $\Theta(n)$  by master theorem

- ▶ worst case is the same for quicksort, extreme pivot at each step,  $\Theta(n^2)$  time
- ▶ **takeaway:** randomized selection takes  $\Theta(n)$  expected time and  $\Theta(n^2)$  worst-case time

## Deterministic Selection Overview

- ▶ **deterministic:** perfectly predictable; not randomized
- ▶ recall that  $T(n) = T(fn) + \Theta(n)$  is  $\Theta(n)$  for any fraction  $0 < f < 1$ , not just  $f = 1/2$
- ▶ need: deterministic process to find a not-terrible pivot
- ▶ i.e. need at least  $fn$  elements on each side of the pivot, so that the worst-case recursive call is  $T((1 - f)n)$
- ▶ e.g. need at least  $\frac{1}{3}n$  elements on each side of the pivot, so that there is a  $T(\frac{1}{3}n)$  or  $T(\frac{2}{3}n)$  call; worst-case is  $T(\frac{2}{3}n)$ ; so

$$T(n) = T\left(\frac{2}{3}n\right) + \Theta(n)$$

which is still  $\Theta(n)$  (though with worse constants)

## Deterministic Selection Process

1. divide  $n$  elements into  $\approx n/5$  groups of 5 elements each
2. find the median of each group with  
*SELECTION – BY – SORTING*;  $\Theta(n(5 \log 5)) = \Theta(n)$  time
3. form a new array of the medians, and recursively select the median of this array = “median-of-medians”;  $T(n/5)$  time
4. partition as usual, using median-of-medians as the pivot;  $\Theta(n)$  time
5. same three cases: either pivot is answer, or recurse before pivot, or recurse after pivot;  
 $T(\text{max. \# elements on either side of pivot})$



## Deterministic Selection Analysis

- ▶ let  $x$  be the median-of-medians; count elements  $\geq x$
- ▶ suppose W.L.O.G. that input elements are distinct
- ▶  $\therefore$  at least half of the group-medians are  $\geq x$
- ▶  $\therefore$  at least half of the groups contain at least 3 elements  $\geq x$  each; except for the group containing  $x$ , and possibly one group with  $< 5$  elements
- ▶  $\therefore$  #elements  $\geq x$  is at least

$$3\left(\left\lceil \frac{1}{2} \left\lceil \frac{n}{5} \right\rceil \right\rceil - 2\right) \geq \frac{3}{10}n - 6$$

- ▶ symmetrically there are at least  $\frac{3}{10}n - 6$  elements  $\leq x$
- ▶  $\therefore$  recursively select at most  $n - (\frac{3}{10}n - 6) = \frac{7}{10}n + 6$  elements

## Deterministic Selection Analysis (continued)

For some  $t \in \Theta(1)$ ,

$$T(n) \leq \begin{cases} O(1) & n < t \\ T(\lceil n/5 \rceil) + T(\frac{7}{10}n + 6) + O(n) & n \geq t. \end{cases}$$

The master theorem does not apply, but the substitution method can be used to show  $T(n) \in O(n)$ .

**Takeaway:** Deterministic selection takes  $O(n)$  worst-case time.

Surprise: selection can be **derandomized** from  $O(n)$  expected time to  $O(n)$  worst-case time with no asymptotic overhead.

Impractical; much worse constant factors, not usually worth it.