

07. Amortized Analysis and Fibonacci Heaps

CPSC 535 ~ Spring 2019

Kevin A. Wortman



CALIFORNIA STATE UNIVERSITY
FULLERTON

March 11, 2019



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

Amortized Efficiency

amortized performance

- ▶ *amortized efficiency*: average of worst-case time over any sequence of operations
- ▶ different from (probabalistic) expected efficiency, which is averaged over random choices made by a PRNG
- ▶ amortized time bound is a weaker math condition than the same worst-case bound
- ▶ e.g. $\Theta(1)$ amortized vs. $\Theta(1)$ worst-case
- ▶ \implies downgrading to an amortized bound may admit upgrades to other aspects of a data structure

Big Idea in Algorithm Design: Work Smart not Hard

Lifecycle of an implemented algorithm:

Phase of Life

design and analysis

learned by students

implemented

run and maintained

Frequency of Creation

once, by discoverer and peer reviewers

annually, by thousands

once per programming language lifetime (decade)

millions-billions of times indefinitely

- ▶ given the choice, we'd prefer for the ongoing tasks to be easy
- ▶ even if that means the one-time phases are complicated
- ▶ holy grail: algorithm is tough to conceive and analyze, but easy to understand and implement
- ▶ example: universal hashing, open addressing

Amortized Analysis — What to Prove

A *splay tree* is a kind of binary search tree.

Lemma: the *INSERT*, *SEARCH*, and *DELETE* operations each take $\Theta(\log n)$ amortized time on a splay tree.

Would need to prove:

- ▶ average time/operation = $\Theta(\log n)$
- ▶ or, any sequence of n of these operations takes a grand total of $n \cdot \Theta(\log n) = \Theta(n \log n)$ worst-case time
- ▶ any sequence: includes the worst-case for the data structure
- ▶ three conventional proof techniques

Aggregate Analysis

- ▶ count total time $T(n)$ for any sequence of n operations
- ▶ \implies each operation takes $T(n)/n$ amortized time
- ▶ pro: simple logic; from first principles
- ▶ con: only works when the goal is to prove all operations take the same time (e.g. *INSERT*, *SEARCH*, and *DELETE* are each $\Theta(\log n)$); wouldn't work if one of them is $\Theta(1)$
- ▶ con: not much inspiration for difficult analyses

The Accounting Method

- ▶ inspired by financial amortization, balanced budgets
- ▶ analyst picks a **cost** for each operation (ex.: *INSERT* costs $4 \cdot \lceil \log_2 n \rceil$ units)
- ▶ some operations **over-charge** and turn a **profit** that is deposited somewhere in the structure (e.g. in a node)
- ▶ other operations **under-charge** and incur a **loss** that is withdrawn from prior deposits
- ▶ show: never go bankrupt, i.e. withdrawn units always exist
- ▶ show: for every operation,

$$(\text{charge}) + (\text{withdrawal}) \leq \text{actual time spent}$$

- ▶ \implies amortized time = charged cost

Accounting Method Pros/Cons

- ▶ pro: possible to prove different amortized efficiency classes for each operation (e.g. one is $\Theta(1)$, another $\Theta(\log n)$)
- ▶ pro: cost story helps us reason through analysis
- ▶ loss operation = procrastinating, deferring work to later
- ▶ profit operation = catching up on deferred work
- ▶ con: cost story may overcomplicate things
- ▶ con: sometimes awkward to store profits in specific data structure locations

Potential Method

- ▶ premise: for data structure D , **potential function** $\Phi(D)$
- ▶ (Φ pronounced like “fee”)
- ▶ potential is stored energy, like a battery
- ▶ amortized time of operation = actual time + change to $\Phi(D)$
- ▶ procrastinating/loss operations decrease $\Phi(D)$
- ▶ catch-up/profit operations increase $\Phi(D)$
- ▶ show: $\Phi(D) \geq 0$ always
- ▶ ultimately a less structured way of thinking of the accounting method

Fibonacci Heap Overview

- ▶ alternative to a binary heap (as in heapsort)
- ▶ pro: faster INSERT and DECREASE-KEY operations (optimal in fact)
- ▶ pro: supports additional operations
- ▶ con: more complicated to describe, implement, *especially* analyze
- ▶ con: operations have amortized time bounds
- ▶ con: much worse constant factors, not in-place
- ▶ not currently practical

Mergeable Heap API

CREATE-HEAP(): initialize empty heap

INSERT(H, x): insert entry x into H

MINIMUM(H): return minimum-key entry

EXTRACT-MIN(H): remove and return minimum-key entry

UNION(H_1, H_2): consume heaps H_1, H_2 and return a new heap with their elements

DECREASE-KEY(H, x, k): lower key of x to k

DELETE(H, x): remove entry x from H

Heap Efficiency Comparison

Operation	Binary Heap (worst-case)	Fib. Heap (amort.)
CREATE-HEAP	$\Theta(1)$	$\Theta(1)$
INSERT	$\Theta(\log n)$	$\Theta(1)$
MINIMUM	$\Theta(1)$	$\Theta(1)$
EXTRACT-MIN	$\Theta(\log n)$	$O(\log n)$
UNION	$\Theta(n)$	$\Theta(1)$
DECREASE-KEY	$\Theta(\log n)$	$\Theta(1)$
DELETE	$\Theta(\log n)$	$\Theta(\log n)$

Fibonacci Heap Structure

Overall heap H has

- ▶ a *forest* of individual k -ary trees (nonbinary); each tree in min-heap-order
- ▶ $H.min$ = pointer to the root with the global minimum key
- ▶ $H.n$ = number of entries in heap

Each node x has

- ▶ $x.parent$ = parent node (NIL if root)
- ▶ $x.child$ = pointer to an arbitrary child
- ▶ $x.left, x.right$ = pointers to siblings
- ▶ (siblings are in circular, unsorted, doubly-linked lists)
- ▶ $x.degree$ = # children
- ▶ $x.mark$ = boolean, true iff x has lost a child since the last time x was assigned a new parent (manages procrastination)

Potential Function and Degree Bound

For analysis purposes, we define the potential function

$$\Phi(H) = t(H) + 2 \cdot m(H)$$

where

- ▶ $t(H) = \#$ trees in root list
- ▶ $m(H) = \#$ marked nodes

Claim: If Fibonacci heap H has n nodes, then the maximum degree ($\#$ children) of any node, written $D(n)$, is

$$D(n) = O(\log n).$$

CREATE-HEAP

```
1: function CREATE-HEAP
2:    $H.min = \text{NIL}$ 
3:    $H.n = 0$ 
4:   return  $H$ 
5: end function
```

Takes $\Theta(1)$ time; $t(H) = m(H) = 0$, so $\Phi(H) = 0$; no profit or loss.

INSERT Pseudocode

Just create node, initialize node, insert into root list, update $H.min$

```
1: function INSERT( $H, x$ )
2:    $x.degree = 0$ 
3:    $x.parent = x.child = \text{NIL}$ 
4:    $x.mark = \text{false}$ 
5:   if  $H.min == \text{NIL}$  then
6:      $H.min = x$ 
7:      $x.left = x.right = \text{NIL}$ 
8:   else
9:     insert  $x$  into root list as  $H.min$ 's right sibling
10:    if  $x.key < H.min.key$  then
11:       $H.min = x$ 
12:    end if
13:  end if
14:   $H.n++$ 
15: end function
```

INSERT Analysis

Actual time steps are $\Theta(1)$.

Let H be heap before INSERT, and H' after; then

$$t(H') = t(H) + 1 \text{ (one tree created)}$$

$$m(H') = m(H) \text{ (no marking)}$$

$$\implies \Phi(H') = \Phi(H) + 1$$

\therefore INSERT takes $O(1) + 1 = O(1)$ amortized time.

MINIMUM

Just follow the $H.min$ pointer.

$O(1)$ actual time.

No new trees, no new marked nodes, so $O(1)$ amortized time.

UNION

```
1: function UNION( $H_1, H_2$ )
2:   if  $H_1.n == 0$  then
3:     return  $H_2$ 
4:   else if  $H_2.n == 0$  then
5:     return  $H_1$ 
6:   else
7:      $H =$  new heap object
8:      $H.min = \min(H_1.min, H_2.min)$ 
9:      $H.n = H_1.n + H_2.n$ 
10:    concatenate root lists of  $H_1, H_2$  into one linked list
11:    return  $H$ 
12:  end if
13: end function
```

$O(1)$ actual time; trees and marked nodes move around, but their number is unchanged, so no change to Φ ; $O(1)$ amortized time.

So Far So Good

So far all operations have been simple, and either potential-neutral (CREATE-HEAP, MINIMUM), or increased potential (INSERT)

Indeed, n INSERTs creates a glorified n -element linked list.

∴ need to expect remaining operations to be more complicated, decrease potential, decrease length of root list.

EXTRACT-MIN

Follow $H.min$ pointer; promote all children to roots; remove from root list; CONSOLIDATE to shorten root list and find new $H.min$.

```
1: function EXTRACT-MIN( $H$ )
2:    $z = H.min$ 
3:   for each child  $c$  of  $z$  do
4:     insert  $c$  into  $H$ 's root list
5:      $c.parent = NIL$ 
6:   end for
7:   delete  $z$  from root list
8:   if  $H.n == 1$  then
9:      $H.min = NIL$                                      ▷ heap just became empty
10:  else
11:    CONSOLIDATE( $H$ )                                     ▷ compact root list, recompute  $H.min$ 
12:  end if
13:   $H.n \leftarrow H.n - 1$ 
14:  return  $z$ 
15: end function
```

Zooming in to CONSOLIDATE

Before moving on, observe EXTRACT-MIN takes, aside from CONSOLIDATE, $\Theta(\text{degree}(H.min))$ time; claimed this is $\Theta(\log n)$

Contract for CONSOLIDATE:

1: **function** CONSOLIDATE(H)

Require: $H.min$ invalid and $H.n == \#nodes + 1$

Ensure: $H.min$ valid and $\#trees \in O(\log n)$

2: **end function**

CONSOLIDATE Pseudocode

```
1: function CONSOLIDATE( $H$ )  
2:    $A = \text{UNIQUE-DEGREE-ARRAY}(H)$   
3:    $\text{ARRAY-TO-ROOT-LIST}(H, A)$   
4:   free  $A$   
5:    $\text{RECOMPUTE-MIN}(H)$   
6: end function
```

Clearly $\Theta(1)$ time except for the three subroutines.

UNIQUE-DEGREE-ARRAY Pseudocode

```

1: function UNIQUE-DEGREE-ARRAY( $H$ )
Ensure: returns  $A[0..D(H.n)]$  where  $A[d]$  = only root w/ degree  $d$ 
2:    $A[0..D(H.n)]$  = new array of node pointers, all NIL
3:   for each root node  $r$  in  $H$  do                                ▷ move  $r$  into  $A$ 
4:      $p = r$                                                         ▷ parent node that needs to move into  $A$ 
5:     while  $A[p.degree] \neq \text{NIL}$  do                                ▷ another node in the way
6:        $c = A[p.degree]$                                             ▷  $p, c$  have same degree
7:        $A[p.degree] = \text{NIL}$                                         ▷ we will link them into one tree
8:       if  $p.key > c.key$  then
9:          $\text{swap}(p, c)$                                             ▷ ensure  $c$  should be  $p$ 's child
10:      end if
11:      make  $c$  a child of  $p$ , incrementing  $p.degree$ 
12:       $c.mark = \text{false}$ 
13:    end while
14:  end for
15:  return  $A$ 
16: end function

```

UNIQUE-DEGREE-ARRAY Analysis

Create A : $O(D(n))$ time

for loop — aggregate analysis

- ▶ # iterations = length of root list before CONSOLIDATE = $t(H) + D(n) - 1$
- ▶ each iteration of inner **while** loop links two roots into one, decrementing # roots
- ▶ # roots at end of CONSOLIDATE \leq size of $A = D(n) + 1$
- ▶ \implies total time in all **while** iterations is

$$(t(H) + D(n) - 1) - (D(n) + 1) = t(H) - 2$$

- ▶ \implies total time in **for** loop is

$$(t(H) + D(n) - 1) + (t(H) - 2) \leq 2 \cdot t(H) + D(n)$$

total for UNIQUE-DEGREE-ARRAY is $O(t(H) + D(n))$

ARRAY-TO-ROOT-LIST

```
1: function ARRAY-TO-ROOT-LIST( $H, A$ )
2:   clear  $H$ 's root list to empty
3:   for index  $i$  of  $A$  do
4:     if  $A[i] \neq \text{NIL}$  then
5:       insert  $A[i]$  into  $H$ 's root list
6:     end if
7:   end for
8: end function
```

$O(D(n))$ time

RECOMPUTE-MIN

```
1: function RECOMPUTE-MIN( $H$ )
2:    $H.min = \text{NIL}$ 
3:   for node  $r$  in  $H$ 's root list do
4:     if  $H.min == \text{NIL}$  then
5:        $H.min = A[i]$ 
6:     else if  $H.min.key > r.key$  then
7:        $H.min = r$ 
8:     end if
9:   end for
10: end function
```

$O(D(n))$ time

CONSOLIDATE Worst-Case Analysis

```

1: function CONSOLIDATE( $H$ )
2:    $A = \text{UNIQUE-DEGREE-ARRAY}(H)$        $\triangleright O(D(n) + t(H))$ 
3:    $\text{ARRAY-TO-ROOT-LIST}(H, A)$            $\triangleright O(D(n))$ 
4:   free  $A$                               $\triangleright O(1)$ 
5:    $\text{RECOMPUTE-MIN}(H)$                    $\triangleright O(D(n))$ 
6: end function

```

Time spent is

$$O(D(n) + t(H) + D(n) + 1 + D(n)) = O(3 \cdot D(n) + t(H)) = O(\log n + t(H)).$$

Want this to be $O(\log n)$; the $t(H)$ overage can be withdrawn from the potential function.

CONSOLIDATE Amortized Analysis

Recall potential function

$$\Phi(H) = t(H) + 2 \cdot m(H)$$

Let H' be H after EXTRACT-MIN

EXTRACT-MIN does not mark any nodes, so $m(H') = m(H)$.

roots decreases: $t(H') = D(n) + 1 \leq t(H)$

$$\Phi(H') = t(H') + 2 \cdot m(H') = (D(n) + 1) + 2 \cdot m(H)$$

$$\Phi(H) - \Phi(H') = t(H) - (D(n) + 1)$$

Amortized cost of EXTRACT-MIN

$$= \Delta\Phi = O(t(H) + D(n)) - O(t(H) - D(n) - 1) = O(2 \cdot D(n)) = O(\log n)$$

Where did that time come from? Accounting Perspective

- ▶ potential function's $t(H)$ term over-charges INSERT and under-charges EXTRACT-MIN
- ▶ when $t(H) = O(\log n)$, EXTRACT-MIN takes $O(\log n)$ worst-case time
- ▶ first $\log n$ roots are “clean”, remaining $t(H) - \log n$ are “mess”
- ▶ deposit $O(1)$ time when INSERTing a messy root
- ▶ withdraw those deposits in CONSOLIDATE
- ▶ 1 deposit pays for making 1 messy root a child of a clean parent
- ▶ (recall that link operation takes $O(1)$ time)
- ▶ hard to use accounting method directly because the identity of the clean nodes changes during CONSOLIDATE based on heap-order

DELETE

```
1: function DELETE( $H, x$ )  
2:   DECREASE-KEY( $H, x, -\infty$ )  
3:   EXTRACT-MIN( $H$ )  
4: end function
```

Know EXTRACT-MIN is $O(\log n)$ amortized time.

Claim: DECREASE-KEY is $O(1)$ amortized time.

\implies DELETE is $O(\log n)$ amortized time.

DECREASE-KEY Sketch

DECREASE-KEY(H, x, k)

- ▶ update $x.key = k$
- ▶ if x is a root, or heap-order maintained, done
- ▶ else **cut** x
 - ▶ make x a root, no longer a child of parent p
 - ▶ update $H.min$ if x is new minimum
 - ▶ mark p ; but if p was already marked, recursively cut p
- ▶ $O(1)$ time not counting recursive cuts
- ▶ recursive cuts are paid for by withdrawing potential

Heap Efficiency Comparison

Operation	Binary Heap (worst-case)	Fib. Heap (amort.)
CREATE-HEAP	$\Theta(1)$	$\Theta(1)$
INSERT	$\Theta(\log n)$	$\Theta(1)$
MINIMUM	$\Theta(1)$	$\Theta(1)$
EXTRACT-MIN	$\Theta(\log n)$	$O(\log n)$
UNION	$\Theta(n)$	$\Theta(1)$
DECREASE-KEY	$\Theta(\log n)$	$\Theta(1)$
DELETE	$\Theta(\log n)$	$\Theta(\log n)$