

05. Hash Tables

CPSC 535 ~ Spring 2019

Kevin A. Wortman



February 25, 2019



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

Big Idea in Algorithm Design

Recall that radix sort takes $\Theta(n)$ time

- ▶ linear time is extremely fast
- ▶ same asymptotically as merely observing all the elements
- ▶ under the assumptions of radix sort, sorting is “free”; no significant computation time

Natural questions:

- ▶ How far can we take this?
- ▶ What is the most powerful way we can organize data, and still only spend $O(1)$ time per element?
- ▶ What are the trade-offs (assumptions) and are they worth it?

The Story So Far

- ▶ *offline*: all operations performed at once; e.g. BUILD-HEAP
- ▶ *online*: data structure valid after each individual operation; e.g. DELETE-MIN
- ▶ computing order offline = sorting problem = $\Omega(n \log n)$ lower bound = $\Omega(\log n)$ /element
- ▶ maintaining order online = search tree = $\Theta(\log n)$ /element
- ▶ sorting integers = radix sort = $\Theta(n)$ total = $\Theta(1)$ /element; subjectively more complex
- ▶ Question: what about maintaining order of integers online, $\Theta(1)$ per element?

Introducing Hash Tables

- ▶ (later) van Emde Boas trees can maintain order of integers in $\Theta(\log \log n)$ per element (huh?)
- ▶ (today) hash tables can't maintain order, but can support dictionary operations INSERT/SEARCH/DELETE, in $\Theta(1)$ /element **expected** time
- ▶ hash tables match runtime of arrays, but (again) are more complicated to understand and implement
- ▶ success story for algorithm design
- ▶ mainstream; taken for granted in CS systems design
- ▶ GZIP, DNS, Java, Python, JSON, ...

Dictionary Operations

CREATE(T): initialize T as a valid dictionary

INSERT(T, k, v): associate key k and value v in dictionary T ; k must not already be in T

SEARCH(T, k): return the value v associated with key k in dictionary T ; or return NIL if k is absent from T

DELETE(T, k): remove key k and its associated value from dictionary T ; k must already be in T

(Practical implementations are often more flexible about re-insertions and ineffectual deletions.)

Direct Address Table

Suppose the *universe* of keys is $k \in U = \{0, 1, \dots, m-1\}$ for fixed m ; create m -element array; $T[k] = v$ (or NIL if k is absent)

```
1: function DA-CREATE(T)
2:    $T[0 \dots m-1] = \text{NIL}$ 
3: end function
```

```
1: function DA-INSERT(T, k, v)
2:    $T[k] = v$ 
3: end function
```

```
1: function DA-SEARCH(T, k)
2:   return  $T[k]$ 
3: end function
```

```
1: function DA-DELETE(T, k)
2:    $T[k] = \text{NIL}$ 
3: end function
```

$\Theta(m)$ space; CREATE is $\Theta(m)$ time; rest are $\Theta(1)$ time; good when $m \in O(n)$ but m could be exponential in n

Hash Functions

- ▶ U = set of keys, m = size of array
- ▶ $m = |U|$ impractical in general
- ▶ introduce *hash function* $h : U \mapsto \{0, 1, \dots, m - 1\}$
- ▶ h “compresses” large space of keys U into denser space of table indices $\leq m$
- ▶ *collisions* exist, i.e. $k_1, k_2 \in U$, $k_1 \neq k_2$, such that

$$h(k_1) = h(k_2)$$

(pigeonhole principle)

- ▶ hard part is designing a concrete h , and collision resolution strategy

Collision Resolution by Chaining

Chaining: each $T[i]$ is a linked list of (colliding) key-value entries

1: **function**

 CHAINED-CREATE(T)

2: $T[0 \dots m - 1] = \text{empty list}$

3: **end function**

1: **function**

 CHAINED-INSERT(T, k, v)

2: insert (k, v) at front of
 $T[h(k)]$

3: **end function**

1: **function**

 CHAINED-SEARCH(T, k)

2: search for k in $T[h(k)]$

3: **end function**

1: **function** DA-DELETE(T, k)

2: remove k from $T[h(k)]$

3: **end function**

CREATE is $O(m)$; INSERT is $O(1)$; SEARCH and DELETE are $O(|T[h(k)]|)$

Chaining Analysis

SEARCH, DELETE take $\Theta(n)$ time in worst case

Let $\alpha = n/m = \text{average keys/table element} = \mathbf{\text{load factor}}$

simple uniform hashing assumption: any element is equally likely to hash to each index, independent of other elements

Under this assumption, expected chain length is $O(1 + \alpha)$
 \implies SEARCH, DELETE take $O(1 + \alpha)$ expected time each

To achieve $O(1 + \alpha) = O(1)$ exp. time, choose $m \in \Omega(n)$

e.g. choose $m \geq 2n$ so

$$\alpha = \frac{n}{m} \leq \frac{n}{(2n)} = \frac{1}{2}$$

Deterministic Hash Functions

First: map keys to natural numbers $\{0, 1, \dots\}$;

henceforth assume $k \in \mathbb{N}$

Division method: $h(k) = k \bmod m$,

for appropriate m

e.g. m is a prime “not too close to” a power of 2

Multiplication method: $h(k) = \lfloor m(kA - \lfloor kA \rfloor) \rfloor$

for carefully-chosen $0 < A < 1$

e.g. $A \approx (\sqrt{5} - 1)/2$

Each is fast, but deterministic, so an adversary could craft a set of keys that force $\Theta(n)/\text{element} \implies \Theta(n^2)$ to insert n elements

Hash Collision Exploits

Denial-of-service attacks based on hash collisions are a real thing

- ▶ CVE-2011-4885
- ▶ working exploit

Universal Hashing

Goal: prevent hash collision exploit

Hash function not compiled in; picked randomly at runtime (once per boot, execution, or table creation); unknowable to adversary

A set \mathcal{H} of hash functions is **universal** when

- ▶ for any distinct keys $k, l \in U$,
- ▶ the number of hash functions for which k and l collide, is at most

$$\frac{|\mathcal{H}|}{m}$$

- ▶ i.e. if h is picked randomly from \mathcal{H} , then for **any** keys k, l ,

$$Pr[h(k) = h(l)] = Pr[\text{two random table indices match}] = \frac{1}{m}$$

Universal Hashing Analysis

Theorem: If hash function h is chosen randomly from a universal family; T is a chained hash table using h , with m table elements, containing n keys, and with load factor $\alpha = n/m$; k is a query key that hashes to list ℓ ; then the expected length of ℓ is α if $k \in T$, or $(1 + \alpha)$ if $k \notin T$.

Note: the hash function type and collision resolution algorithm need to be analyzed together.

Corollary: In a hash table using a universal hash function and chaining collision resolution, INSERT takes $O(1)$ time, while SEARCH and DELETE each take $O(1)$ expected time and $\Theta(n)$ worst-case time.

An Example Universal Hash Family

Fix p as a prime number greater than every possible key k

Let a, b be int. parameters with $1 \leq a < p$ and $0 \leq b < p$; define

$$h_{ab}(k) \equiv ((ak + b) \bmod p) \bmod m.$$

Number theory shows this is a universal family.

Conveniences

- ▶ can hardcode p as a prime near $2^W - 1$
- ▶ generating h_{ab} only involves choosing two random int's $< p$
- ▶ evaluating h_{ab} only involves one integer multiply, one add, and two modulo operations
- ▶ can recalibrate a fixed h_{ab} to any m

Open Addressing

Concern with chaining: linked lists have poor *locality of reference*, require expensive allocate/free operations, and are conceptually complicated.

open addressing: all key-value entries are stored directly in the table T itself; there are no chains

Modest assumptions

- ▶ each $T[i]$ can store either an active (k, v) entry; or NIL (meaning never used); or TOMBSTONE (meaning was once active, but is now empty)
- ▶ $n < m$, i.e. $\alpha < 1$

Open Addressing Delete

DELETE:

- ▶ let $i = h(k)$; **probe** (examine) $T[i]$
- ▶ if k is found, $T[i] = \text{TOMBSTONE}$, done
- ▶ else if $T[i]$ is NIL, conclude $k \notin T$, done
- ▶ else ($T[i]$ is TOMBSTONE or key other than k), probe a different index
- ▶ after m probes, conclude $k \notin T$, stop (otherwise infinite loop)

probe sequence: sequence of indices probed; different approaches

$O(|\text{probe sequence}|)$ time

Open Addressing Search and Insert

SEARCH:

- ▶ probe in the same order as DELETE
- ▶ keep searching past TOMBSTONE elements

INSERT:

- ▶ probe in the same order as DELETE
- ▶ overwrite the first $T[i]$ containing NIL or TOMBSTONE

Probe Sequence Functions; Linear Probing

Hash function: $h : U \mapsto \{0, 1, \dots, m - 1\}$

Probe sequence function:

$h : U \times \{0, 1, \dots, m - 1\} \mapsto \{0, 1, \dots, m - 1\}$, i.e.

$h(k, i) = i$ th index to probe for key k

Linear probing: probe in sequential order, so

$$ps(k, i) = (h(k) + i) \bmod m$$

Simple; excellent locality of reference; but develops long **clusters** of occupied elements; an empty slot after ℓ occupied slots will be filled with probability $(\ell + 1)/m$

Quadratic Probing

Quadratic probing:

$$ps(k, i) = (h(k) + c_1i + c_2i^2) \mod m$$

for constant parameters c_1, c_2 (with some constraints)

Much less clustering than linear probing; but still some, because if k_1, k_2 collide because $h(k_1) = h(k_2)$, then their entire probe sequences collide.

\implies partially mitigates the clustering problem; but the ps function is more expensive, and locality of reference is not as good

Double Hashing

Double hashing:

$$ps(k, i) = (h_1(k) + i \cdot h_2(k)) \bmod m$$

where h_1, h_2 are two distinct hash functions

Since h_1, h_2 are different, even if k_1, k_2 collide at first with $h_1(k_1) = h_1(k_2)$, their probe sequences are almost certainly different

\implies clustering problem essentially solved; but more expensive than linear/quadratic programming because we evaluate two hash functions, and locality of reference is poor (random), though still entirely within T

Open Addressing Analysis

Universal hashing assumption: the probability of every possible probe sequence $\langle p_1, p_2, \dots, p_m \rangle$ with each $p_i \in \{0, 1, \dots, m-1\}$ is equally likely.

Theorem: Under the universal hashing assumption, an open-address hash table with load factor $\alpha < 1$ performing a search makes at most

$$\frac{1}{\alpha} \ln \frac{1}{1 - \alpha}$$

probes when the search succeeds, or at most

$$\frac{1}{1 - \alpha}$$

probes when the search fails.

Corollary: Inserting requires at most $\frac{1}{1 - \alpha}$ probes on average.

Corollary: If $\alpha \in O(1)$ then the INSERT, SEARCH, and DELETE operations on an open-addressing hash table each take $O(1)$

Perfect hashing

Static hashing: set of keys is known statically (algorithm design-time or program compile-time)

- ▶ CREATE, INSERT happen offline (all-at-once) at compile-time
- ▶ SEARCH happens online at run-time
- ▶ DELETE not supported

Perfect hashing: static hashing where SEARCH $O(1)$ worst case time (**not** $O(1)$ expected time)

Applies when key set is constant, or evolves very slowly: top-level domains, programming language keywords, state abbreviations (CA), etc.

Two-Level Hashing

Hash-table-of-hash-tables:

- ▶ *outer* hash table w/ function h
- ▶ slot $T[j]$ holds an *inner* table using function h_j
- ▶ all hash functions h, h_j are universal
- ▶ if $T[j]$ holds n_j colliding keys, then its size m_j must be $\Omega(n_j^2)$
- ▶ non-obvious: large n_j^2 size prevents collisions in the inner tables but the whole table still only takes $O(n)$ space

Two-Level Collisions

Theorem: if inner hash table $T[j]$ uses a universal hash function h_j and has $m_j = n_j^2$, then the probability of a collision in $T[j]$ is less than $\frac{1}{2}$

Sketch: there are $\binom{n_j}{2}$ pairs of keys; due to universal h_j , each collides with probability $\frac{1}{m}$; so

$$E[\# \text{ collisions}] = \binom{n_j}{2} \cdot \frac{1}{m_j} = \binom{n_j}{2} \cdot \left(\frac{1}{n_j^2}\right) = \left(\frac{n_j^2 - n}{2}\right) \cdot \frac{1}{n_j^2} < \frac{1}{2}.$$

To pick an h_j , keep retrying until no collisions $\implies O(1)$ expected trials

Two-Level Analysis

Theorem: total space used by a two-level hash is $O(n)$

Sketch:

- ▶ outer hash uses $O(n)$ space
- ▶ each $T[j]$ has n_j elements and uses $\Theta(n_j^2)$; space
- ▶ $O(1)$ collisions in $T[j] \implies n_j \in O(1) \implies n_j^2 \in O(1)$

Takeaway: two-level perfect hash table takes $O(n)$ expected time to create statically; a SEARCH takes $O(1)$ deterministic worst-case time at runtime.